

THIS DOCUMENT CONTAINS QUESTIONS THAT REPRESENT THE SORT OF QUESTIONS THAT MIGHT APPEAR ON THE FINAL QUIZ FOR DATA MINING FOR BUSINESS ANALYTICS.

THESE ARE INTENDED TO REPRESENT THE FORMAT AND STYLE OF QUESTIONS, NOT NECESSARILY THE CONTENT.

Chapter-spanning Questions

- 1) Which data science method is most appropriate for the following business question?
"Of all my accounts, which are the most likely to exhibit fraud, based on my experience with prior cases of accounts that have and have not been defrauded?"
 - a) Classification tree induction
 - b) Hierarchical clustering
 - c) k-Means
 - d) Linear regression

- 2) Which analytics technology would be most useful in answering the following business question? "If this customer responds to my offer, how much will she spend?"
 - a) Classification tree induction
 - b) Hierarchical clustering
 - c) k-Means
 - d) Linear regression

- 3) I'm proposing a foreclosure-classification system to a small bank with stakeholders of very varied backgrounds. For a pilot study demonstration on a small data set, which technique is most suitable?

- a) Tree induction
 - b) k-Nearest Neighbor
 - c) Logistic Regression
- 4) Tree induction and clustering both can be used to segment customers. Contrast the process for using these two different types of modeling for customer segmentation. For what sort of problems would you use each? What would you have to do differently?
- 5) I want to rank credit applicants by their estimated likelihood of default. Which technique would be least helpful in assessing the quality of a ranking model mined from data?
- a) holdout testing
 - b) calculate area under the ROC curve
 - c) calculate percent correctly classified instances
 - d) cross-validation
 - e) domain knowledge validation
- 6) I have a business application where I want to make fast classifications, and update my model quickly and immediately when a new training case comes along. What induction algorithm is best suited?
- a) Naïve Bayes
 - b) Classification tree induction
 - c) k-Nearest Neighbor
 - d) Logistic Regression

Matching: Choose the *best* matching for each set; each letter should be used once (these might be hard)

____ cross-validation ____ domain-knowledge validation ____ ROCcurve ____ overfitting avoidance	a. Ranking b. Comprehensibility c. Generalization performance d. Complexity control
--	--

____ holdout evaluation ____ domain-knowledge validation ____ learning curve ____ overfitting ____ increasing comprehensibility	a. Sanity checking b. Pruning c. Increasing data d. Cross-validation e. Divergence between testing accuracies
---	---

____ learning curves ____ fitting curves for kNN ____ information gain ____ cumulative response curves ____ increasing Mann-Whitney-Wilcoxon	a. increasing proportion targeted b. increasing tree size c. increasing training data d. increasing complexity e. increasing AUC
--	--

____ entropy ____ logistic ____ information gain ____ accuracy ____ regression ____ lift	a. log odds b. numeric target c. how mixed up classes are d. higher on training data e. difference between parents and children f. better with model than without
---	--

Free-response Questions

Q) Blue Moon Consulting

After a few beers your CIO invited his buddy from Blue Moon consulting to propose a project using data mining to improve the targeting of the new service that you have been a principal in developing. The service has been quite successful so far, being marketed over the last 6 months via your ingenious, and very inexpensive, word-of-mouth campaign. You've already garnered a pretty large customer base without any targeting, and you've been seeing this success as your best stepping stone to bigger and better things in the firm. After some reflection, you've decided that your best course of action is to play a key role in ensuring the success of the data mining project as well. You agree with your CIO's statement in a meeting with Blue Moon, that very accurate targeting might cost-effectively expand your audience to consumers that word-of-mouth would not reach.

Identify the four most serious flaws in this abridged version of Blue Moon's proposal, and suggest how to ameliorate them. You can accept that Blue Moon has accurate information about the service.

"We will build a logistic regression (LR) model to predict service uptake for a consumer, based on the data on your existing customers, including their demographics and their usage of the service. We believe that logistic regression is the best choice of method because it is a tried and true statistical technique, and we can interpret the coefficients of the model to infer whether the attributes are statistically significant, and whether they make sense. If they are and they do, then we can have confidence that the model will be accurate in predicting service uptake. We will apply the model to our (Blue Moon's) large database of consumers, and select out those whom the LR model predicts to be the most likely to subscribe. It is a fixed-price, fixed-cost, fixed-term service, so this also will in effect rank them by expected profit as well."

Answer

Q) A Political Campaign

Immediately after the successful 2012 Obama campaign, Campaign Manager and former White House Deputy Chief of Staff for Operations Jim Messina discussed their "Targeted Sharing" analytics program: "I think it's one of the most important things we did." Targeted Sharing (TS) addressed the problem that young voters were important for the Obama campaign, but were difficult to reach via the traditional channels employed by political campaigns (direct mail, phone, email). The TS solution involved targeting the friends of the roughly half-million "authorized" Facebook Obama supporters. The analytics involved prioritizing the friends for the "targeted" sharing of campaign content, as oversharing leads to Facebook deprioritizing the content in the friends' streams.

Your liberal leanings led you to volunteer, and now you have been selected to lead the design of an analytics method for prioritized targeting of the 70 million unique friends for a "getting out the vote" message. Choose one of the methods we have covered this semester. Explain your design by answering the following questions. The notion of "proxies" for unavailable data may be helpful. One or two sentences will suffice to answer each question.

- 1) What precisely is your problem formulation? What general category of data mining task does this correspond to?
- 2) What is your data representation? Is this a supervised or unsupervised formulation? Why? What features will you use? Describe a few elements of your feature vector precisely.
- 3) What method do you propose to use? Why?
- 4) How will you evaluate whether your model has captured any generalizable knowledge? Explain two different ways, and what metric you propose to employ.
- 5) If the evaluation shows that it has indeed captured generalizable knowledge, why will (help) to solve the prioritization problem? Explain precisely. Are there other important factors that should be included in the solution?

Q) You sell IT products and your lone data scientist proposes to use k-NN to build an IT wallet estimation predictor to predict the "wallet share" that companies spend with you. Wallet share is the percentage of a company's total IT budget that is spent with your company. You have information on the total IT budgets of a large set of companies, which will constitute your database of potential neighbors. The two of you have already gone over the distance measure and have chosen k. Now you want to estimate your wallet share for Acme Corp., one of your current customers for whom you do not know the IT budget.

- a) Explain precisely how you will estimate your wallet share for Acme with this technique.
- b) If you chose $k=N$, the total number of training examples, what would be the effect?

Q) Your data science team is presenting you the results of their latest analysis.

- A) They present learning curves. Briefly compare and contrast learning curves and fitting curves, including what they show and why they are important.
- B) They are asking for a budget to acquire more labeled training data. How can the learning curve help you to decide whether the request is justified?
- C) They discuss results using cross-validation. What is cross-validation, and why would one use it?
- D) In reporting their results, the team reports the area under the ROC Curve. What the heck is an ROC curve? When is it used and what does it show? Why should I care about the area under it?

Q) I would like to see whether my investment customers tend to cluster in understandable groups. What precisely is the difference between segmenting my customers using clustering and segmenting my customers using tree induction, and when would I use one rather than the other? What is the practical difference—i.e., in the data mining process where and what would the main differences be? What different clustering methods should I consider? Once I get clusters, it is important for me to see whether they make sense. Describe three ways to help understand the resultant clustering.

Q) We would like to target some subset of the huge number of visitors to our main retail web page with a new special offer. Instead of the normal early May special offer of a discounted flower bouquet for Mom, we've decided to offer selected customers a 30% discount on any electric razor purchase from our stock. We need to decide which customers will get the normal special offer and which will get the new special offer.

(a) Show how computing expected value provides a framework for thinking about what models need to be built for this problem.

(b) Specify what models you would build.

(c) Do you expect to have the data necessary to build these models? If so, from where, if not, what do you propose to do about it?

Q) Consider our telco churn example from class. We would like to target some subset of our customers with a brand new special offer prior to the expiration of their contracts, in order to minimize our losses due to churn. The special offer has a fixed cost of **C** irrespective of whether the customer accepts it. You might recall that we already did a similar expected-value based selection of customers in one of our homeworks. In real life however, you might want to design your solution more carefully and differently in some aspects.

(a) How would you approach this problem in reality from the point of view of expected value framework in order to select the subset of customers who should be targeted? Describe carefully your procedure.

(b) Assume you have a certain budget of \$5,000,000 for your campaign. How would you decide which customers to target? How would you determine how many of these customers you will target?

(c) Specify what model or models you would suggest building, including specifying the target variable(s).

(d) Describe the types / groups of attributes you would consider; be as complete as you can while being brief.

(e) Do you expect to have the data necessary to build these models? If so, from where, if not, what do you propose to do about it? You can ignore any value of the customer if she leaves the company. You can assume that the value of a particular customer, if she were to stay, can be estimated directly using this particular customer's past value and knowledge of the business going forward and the incentive, so you don't need to use data mining to estimate this quantity.

Q) You have built a logistic regression model to help in your decision of which customers to target with a special offer prior to contract expiration. The model estimates the probability that a customer will leave within 90 days of contract expiration. You realize however that some customers are more valuable to you than others, and you can acquire data to quantify this. You have a budget of \$10,000 to target customers with a special offer. Explain how you will use the results of your logistic regression model to target your customers.